

# Reinforcement Learning

CS 760@UW-Madison



# Goals for the lecture



you should understand the following concepts

- the reinforcement learning task
- Markov decision process
- value functions
- value iteration
- Q functions
- Q learning
- exploration vs. exploitation tradeoff
- compact representations of Q functions
- reinforcement learning example


# Reinforcement learning (RL)

Task of an agent embedded in an environment

repeat forever

- 1) sense world
- 2) reason
- 3) choose an action to perform
- 4) get feedback (usually reward = 0)
- 5) learn

the environment may be the physical world or an artificial one




# Example: RL Backgammon Player

[Tesauro, *CACM* 1995]



- world
  - 30 pieces, 24 locations
- actions
  - roll dice, e.g. 2, 5
  - move one piece 2
  - move one piece 5
- rewards
  - win, lose
- TD-Gammon 0.0
  - trained against itself (300,000 games)
  - as good as best previous BG computer program (also by Tesauro)
- TD-Gammon 2
  - beat human champion







# Reinforcement learning



- set of states  $S$
- set of actions  $A$
- at each time  $t$ , agent observes state  $s_t \in S$  then chooses action  $a_t \in A$
- then receives reward  $r_t$  and changes to state  $s_{t+1}$



# RL as Markov decision process (MDP)





- Markov assumption

$$P(s_{t+1} | s_t, a_t, s_{t-1}, a_{t-1}, \dots) = P(s_{t+1} | s_t, a_t)$$

- also assume reward is Markovian

$$P(r_{t+1} | s_t, a_t, s_{t-1}, a_{t-1}, \dots) = P(r_{t+1} | s_t, a_t)$$




Goal: learn a policy  $\pi : S \rightarrow A$  for choosing actions that maximizes

$$E[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots] \quad \text{where } 0 \leq \gamma < 1$$

for every possible starting state  $s_0$

# Reinforcement learning task

- Suppose we want to learn a control policy  $\pi : S \rightarrow A$  that maximizes  $\sum_{t=0}^{\infty} \gamma^t E[r_t]$  from every state  $s \in S$



each arrow represents an action  $a$  and the associated number represents deterministic reward  $r(s, a)$



# Value function for a policy

- given a policy  $\pi : S \rightarrow A$  define

$$V^{\pi}(s) = \sum_{t=0}^{\infty} \gamma^t E[r_t]$$

assuming action sequence chosen  
according to  $\pi$  starting at state  $s$

- we want the optimal policy  $\pi^*$  where


$$\pi^* = \arg \max_{\pi} V^{\pi}(s) \quad \text{for all } s$$

we'll denote the value function for this optimal policy as  $V^*(s)$



# Value function for a policy $\pi$

- Suppose  $\pi$  is shown by red arrows,  $\gamma = 0.9$




$V^\pi(s)$  values are shown in red

# Value function for an optimal policy $\pi^*$



- Suppose  $\pi^*$  is shown by red arrows,  $\gamma = 0.9$



$V^*(s)$  values are shown in red

# Using a value function



If we know  $V^*(s)$ ,  $r(s_t, a)$ , and  $P(s_t | s_{t-1}, a_{t-1})$   
we can compute  $\pi^*(s)$

$$\pi^*(s_t) = \arg \max_{a \in A} \left[ r(s_t, a) + \gamma \sum_{s \in S} P(s_{t+1} = s | s_t, a) V^*(s) \right]$$

# Value iteration for learning $V^*(s)$



```
initialize  $V(s)$  arbitrarily
loop until policy good enough
{
  loop for  $s \in S$ 
  {
    loop for  $a \in A$ 
    {

$$Q(s, a) \leftarrow r(s, a) + \gamma \sum_{s' \in S} P(s' | s, a) V(s')$$

    }

$$V(s) \leftarrow \max_a Q(s, a)$$

  }
}
```

# Value iteration for learning $V^*(s)$



- $V(s)$  converges to  $V^*(s)$
- works even if we randomly traverse environment instead of looping through each state and action methodically
  - but we must visit each state infinitely often
- implication: we can do online learning as an agent roams around its environment
- assumes we have a model of the world: i.e. know  $P(s_t | s_{t-1}, a_{t-1})$
- What if we don't?

# $Q$ learning



define a new function, closely related to  $V^*$

$$V^*(s) \leftarrow E[r(s, \pi^*(s))] + \gamma E_{s'|s, \pi^*(s)}[V^*(s')]$$

$$Q(s, a) \leftarrow E[r(s, a)] + \gamma E_{s'|s, a}[V^*(s')]$$


if agent knows  $Q(s, a)$ , it can choose optimal action without knowing  $P(s' | s, a)$

$$\pi^*(s) \leftarrow \arg \max_a Q(s, a) \quad V^*(s) \leftarrow \max_a Q(s, a)$$


and it can learn  $Q(s, a)$  without knowing  $P(s' | s, a)$




# Q values



$r(s, a)$  (immediate reward) values



$V^*(s)$  values



$Q(s, a)$  values

# $Q$ learning for deterministic worlds



for each  $s, a$  initialize table entry  $\hat{Q}(s, a) \leftarrow 0$

observe current state  $s$

do forever

    select an action  $a$  and execute it

    receive immediate reward  $r$


    observe the new state  $s'$

    update table entry

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a')$$

$$s \leftarrow s'$$

# Updating $Q$



$$\begin{aligned}\hat{Q}(s_1, a_{right}) &\leftarrow r + \gamma \max_{a'} \hat{Q}(s_2, a') \\ &\leftarrow 0 + 0.9 \max \{63, 81, 100\} \\ &\leftarrow 90\end{aligned}$$

# $Q$ learning for *nondeterministic* worlds



for each  $s, a$  initialize table entry  $\hat{Q}(s, a) \leftarrow 0$

observe current state  $s$

do forever

    select an action  $a$  and execute it

    receive immediate reward  $r$

    observe the new state  $s'$

    update table entry

$$\hat{Q}_n(s, a) \leftarrow (1 - \alpha_n) \hat{Q}_{n-1}(s, a) + \alpha_n [r + \gamma \max_{a'} \hat{Q}_{n-1}(s', a')]$$

$s \leftarrow s'$

where  $\alpha_n$  is a parameter dependent on the number of visits to the given  $(s, a)$  pair


$$\alpha_n = \frac{1}{1 + \text{visits}_n(s, a)}$$

# Convergence of $Q$ learning



- $Q$  learning will converge to the correct  $Q$  function
  - in the deterministic case
  - in the nondeterministic case (using the update rule just presented)
- in practice it is likely to take many, many iterations

# $Q$ 's vs. $V$ 's



- Which action do we choose when we're in a given state?
- $V$ 's (model-based)
  - need to have a 'next state' function to generate all possible states
  - choose next state with highest  $V$  value.
- $Q$ 's (model-free)
  - need only know which actions are legal
  - generally choose next state with highest  $Q$  value.





# Exploration vs. Exploitation

- in order to learn about better alternatives, we shouldn't always follow the current policy (**exploitation**)
- sometimes, we should select random actions (**exploration**)
- one way to do this: select actions probabilistically according to:

$$P(a_i | s) = \frac{c^{\hat{Q}(s, a_i)}}{\sum_j c^{\hat{Q}(s, a_j)}}$$

where  $c > 0$  is a constant that determines how strongly selection favors actions with higher  $Q$  values

# $Q$ learning with a table



As described so far, Q learning entails filling in a huge table


		states				
		$s_0$	$s_1$	$s_2$	$\dots$	$s_n$
actions	$a_1$			.		
	$a_2$			.		
	$a_3$	$\dots$		$Q(s_2, a_3)$		
	.					
	.					
	$a_k$					

A table is a very verbose way to represent a function

# Representing $Q$ functions more compactly



We can use some other function representation (e.g. a neural net) to compactly encode a substitute for the big table



each input unit encodes  
a property of the state  
(e.g., a sensor value)

or could have one net  
for each possible action

# Why use a compact $Q$ function?



1. Full  $Q$  table may not fit in memory for realistic problems
2. Can **generalize across states**, thereby speeding up convergence  
i.e. one instance 'fills' many cells in the  $Q$  table

## Notes

1. When generalizing across states, cannot use  $\alpha=1$
2. Convergence proofs only apply to  $Q$  tables
3. Some work on bounding errors caused by using compact representations (e.g. Singh & Yee, *Machine Learning* 1994)

# $Q$ tables vs. $Q$ nets



Given: 100 Boolean-valued features

10 possible actions

Size of  $Q$  table

$10 \times 2^{100}$  entries

Size of  $Q$  net (assume 100 hidden units)

$100 \times 100 + 100 \times 10 = 11,000$  weights

weights between  
inputs and HU's

weights between  
HU's and outputs

# Representing $Q$ functions more compactly



- we can use other regression methods to represent  $Q$  functions
  - $k$ -NN
  - regression trees
  - support vector regression
  - etc.



# $Q$ learning with function approximation




1. measure sensors, sense state  $s_0$
2. predict  $\hat{Q}_n(s_0, a)$  for each action  $a$
3. select action  $a$  to take (with randomization to ensure exploration)
4. apply action  $a$  in the real world
5. sense new state  $s_1$  and immediate reward  $r$
6. calculate action  $a'$  that maximizes  $\hat{Q}_n(s_1, a')$
7. train with new instance

$$\mathbf{x} = s_0$$

$$y \leftarrow (1 - \alpha)\hat{Q}(s_0, a) + \alpha[r + \gamma \max_{a'} \hat{Q}(s_1, a')]$$

*Calculate Q-value you would have put into Q-table, and use it as the training label*

# ML example: reinforcement learning to control an autonomous helicopter



video of Stanford University autonomous helicopter from <http://heli.stanford.edu/>


# Stanford autonomous helicopter



## sensing the helicopter's state

- orientation sensor
  - accelerometer
  - rate gyro
  - magnetometer
- GPS receiver ("2cm accuracy as long as its antenna is pointing towards the sky")
- ground-based cameras

## actions to control the helicopter



# Experimental setup for helicopter



1. Expert pilot demonstrates the airshow several times



2. Learn a reward function based on desired trajectory
3. Learn a dynamics model
4. Find the optimal control policy for learned reward and dynamics model
5. Autonomously fly the airshow



6. Learn an improved dynamics model. Go back to step 4

# Learning dynamics model $P(s_{t+1} \mid s_t, a)$



- state represented by helicopter's

position  $(x, y, z)$

velocity  $(\dot{x}, \dot{y}, \dot{z})$

angular velocity  $(\omega_x, \omega_y, \omega_z)$

- action represented by manipulations of 4 controls

$(u_1, u_2, u_3, u_4)$

- dynamics model predicts accelerations as a function of current state and actions
- accelerations are integrated to compute the predicted next state

# Learning dynamics model $P(s_{t+1} \mid s_t, a)$



dynamics  
model

$$\ddot{x}^b = A_x \dot{x}^b + g_x^b + w_x,$$

$$\ddot{y}^b = A_y \dot{y}^b + g_y^b + D_0 + w_y,$$

$$\ddot{z}^b = A_z \dot{z}^b + g_z^b + C_4 u_4 + D_4 + w_z,$$

$$\dot{\omega}_x^b = B_x \omega_x^b + C_1 u_1 + D_1 + w_{\omega_x},$$

$$\dot{\omega}_y^b = B_y \omega_y^b + C_2 u_2 + D_2 + w_{\omega_y},$$

$$\dot{\omega}_z^b = B_z \omega_z^b + C_3 u_3 + D_3 + w_{\omega_z}.$$

- $A, B, C, D$  represent model parameters
- $g$  represents gravity vector
- $w$ 's are random variables representing noise and unmodeled effects
- linear regression task!



# Learning a desired trajectory

- repeated expert demonstrations are often suboptimal in different ways
- given a set of  $M$  demonstrated trajectories

$$y_j^k = \begin{bmatrix} s_j^k \\ u_j^k \end{bmatrix} \quad \text{for } j = 0, \dots, N-1, k = 0, \dots, M-1$$

action on  $j^{\text{th}}$  step of trajectory  $k$       state on  $j^{\text{th}}$  step of trajectory  $k$

- try to infer the implicit desired trajectory

$$z_t = \begin{bmatrix} s_t^* \\ u_t^* \end{bmatrix} \quad \text{for } t = 0, \dots, H$$

# Learning a desired trajectory



colored lines: demonstrations of two loops  
black line: inferred trajectory

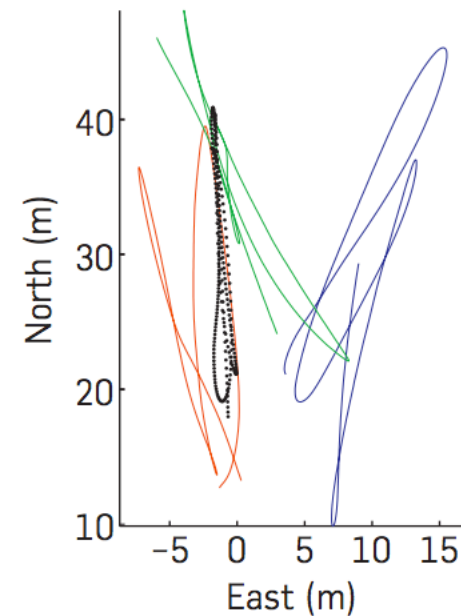



Figure from Coates et al., CACM 2009


# Generative model for desired trajectories



desired trajectory

demonstrated trajectory

time indices



$$z_{t+1} = f(z_t) + w_t^{(z)}, \quad w_t^{(z)} \sim \mathcal{N}(0, \Sigma^{(z)})$$

dynamics model

noise


# Generative model for desired trajectories



desired trajectory

demonstrated trajectory

time indices



$$y_j^k = z_{\tau_j^k} + w_j, \quad w_t^{(y)} \sim \mathcal{N}(0, \Sigma^{(y)})$$

desired trajectory at  
time step indexed by  $\tau_j^k$

noise


# Generative model for desired trajectories



desired trajectory

demonstrated trajectory

time indices



$$P(\tau_{j+1}^k | \tau_j^k) = \begin{cases} d_1^k & \text{if } \tau_{j+1}^k - \tau_j^k = 1 \\ d_2^k & \text{if } \tau_{j+1}^k - \tau_j^k = 2 \\ d_3^k & \text{if } \tau_{j+1}^k - \tau_j^k = 3 \\ 0 & \text{otherwise} \end{cases}$$

parameters specifying probability of the subsampling interval

# Learning reward function



- EM is used to infer desired trajectory from set of demonstrated trajectories
- The reward function is based on deviations from the desired trajectory

# Finding the optimal control policy



- finding the control policy is a reinforcement learning task

$$\pi^* \leftarrow \arg \max_{\pi} E \left[ \sum_t r(s_t, a) \mid \pi \right]$$

- RL learning methods described earlier don't quite apply because state and action spaces are both continuous
- A special type of Markov decision process in which the optimal policy can be found efficiently
  - reward is represented as a linear function of state and action vectors
  - next state is represented as a linear function of current state and action vectors
- They use an iterative approach that finds an approximate solution because the reward function used is quadratic



An aerial photograph of a city harbor at sunset. The sun is low on the horizon, casting a warm, golden glow over the water and the city. Numerous sailboats are scattered across the harbor. The city buildings are visible along the shoreline, and a large hill is in the background.

# THANK YOU

Some of the slides in these lectures have been adapted/borrowed from materials developed by Yingyu Liang, Mark Craven, David Page, Jude Shavlik, Tom Mitchell, Nina Balcan, Elad Hazan, Tom Dietterich, and Pedro Domingos.

