

An aerial photograph of the University of Wisconsin-Madison campus during sunset. The sun is low on the horizon, casting a warm golden glow over the buildings and the surrounding trees. The campus buildings are a mix of architectural styles, with many featuring red brick. A large, calm lake, Lake Mendota, stretches across the right side of the frame. Numerous small sailboats and other watercraft are scattered across the dark blue water. In the foreground, the city skyline is visible, with more buildings and trees. The overall atmosphere is peaceful and scenic.

Instance-Based Learning

CS 760@UW-Madison





Goals for the lecture

you should understand the following concepts

- k -NN classification
- k -NN regression
- edited nearest neighbor
- k-d trees for nearest neighbor identification (optional)
- inductive bias (hypothesis space bias, preference bias)



Nearest-neighbor classification

learning stage

- given a training set $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$, do nothing
(it's sometimes called a *lazy learner*)


classification stage

- **given**: an instance $x^{(q)}$ to classify
- find the training-set instance $x^{(i)}$ that is most similar to $x^{(q)}$
- return the class value $y^{(i)}$



The decision regions

Voronoi diagram: each polyhedron indicates the region of feature space that is in the nearest neighborhood of each training instance





k -nearest-neighbor classification

classification task

- **given:** an instance $x^{(q)}$ to classify
- find the k training-set instances $(x^{(1)}, y^{(1)}), \dots, (x^{(k)}, y^{(k)})$ that are most similar to $x^{(q)}$
- return the class value

$$\hat{y} \leftarrow \operatorname{argmax}_{v \in \text{values}(Y)} \sum_{i=1}^k \delta(v, y^{(i)}) \quad \delta(a, b) = \begin{cases} 1 & \text{if } a = b \\ 0 & \text{otherwise} \end{cases}$$

(i.e. return the class that have the most instances)



How can we determine distance

suppose all features are discrete

- Hamming distance: count the number of features for which two instances differ

suppose all features are continuous

- Euclidean distance:

$$d(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sqrt{\sum_f (x_f^{(i)} - x_f^{(j)})^2} \quad \text{where } x_f^{(i)} \text{ represents the } f\text{-th feature of } \mathbf{x}^{(i)}$$

- Manhattan distance:

$$d(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sum_f |x_f^{(i)} - x_f^{(j)}|$$



How can we determine distance

- if we have a mix of discrete/continuous features:

$$d(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \sum_f \begin{cases} |x_f^{(i)} - x_f^{(j)}| & \text{if } f \text{ is continuous} \\ 1 - \delta(x_f^{(i)}, x_f^{(j)}) & \text{if } f \text{ is discrete} \end{cases}$$

- typically want to apply to continuous features some type of normalization (values range 0 to 1) or standardization (values distributed according to standard normal)
- many other possible distance functions we could use ...



Standardizing numeric features

- given the training set D , determine the mean and stddev for feature x_i

$$\mu_i = \frac{1}{|D|} \sum_{d=1}^{|D|} x_i^{(d)}$$
$$\sigma_i = \sqrt{\frac{1}{|D|} \sum_{d=1}^{|D|} (x_i^{(d)} - \mu_i)^2}$$

- standardize each value of feature x_i as follows

$$\hat{x}_i^{(d)} = \frac{x_i^{(d)} - \mu_i}{\sigma_i}$$

- do the same for test instances, using the same μ_i and σ_i derived from the *training* data

An aerial photograph of a city skyline at sunset. The city is built along a large body of water, with numerous buildings of various architectural styles and heights. The water in the foreground is dark blue, with many small sailboats and other boats scattered across it. The sky is a warm, golden color from the setting sun. The overall scene is a mix of urban architecture and natural beauty.

Variants





k -nearest-neighbor regression

learning stage

- given a training set $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$, do nothing

prediction stage

- **given**: an instance $x^{(q)}$ to make a prediction for
- find the k training-set instances $(x^{(1)}, y^{(1)}), \dots, (x^{(k)}, y^{(k)})$ that are most similar to $x^{(q)}$
- return the value

$$\hat{y} \leftarrow \frac{1}{k} \sum_{i=1}^k y^{(i)}$$



Distance-weighted nearest neighbor

We can have instances contribute to a prediction according to their distance from $x^{(q)}$

classification:

$$\hat{y} \leftarrow \operatorname{argmax}_{v \in \text{values}(Y)} \sum_{i=1}^k w_i \delta(v, y^{(i)}) \quad w_i = \frac{1}{d(x^{(q)}, x^{(i)})^2}$$


regression:

$$\hat{y} \leftarrow \frac{\sum_{i=1}^k w_i y^{(i)}}{\sum_{i=1}^k w_i}$$




Irrelevant features

here's a case in which there is one relevant feature x_1 and a 1-NN rule classifies each instance correctly



consider the effect of an irrelevant feature x_2 on distances and nearest neighbors





Speeding up k -NN

- k -NN is a “lazy” learning algorithm – does virtually nothing at training time
- but classification/prediction time can be costly when the training set is large
- two general strategies for alleviating this weakness
 - don’t retain every training instance (edited nearest neighbor)
 - use a smart data structure to look up nearest neighbors (e.g. a k-d tree)



Edited instance-based learning

- select a subset of the instances that still provide accurate classifications
- *incremental deletion*
 - start with all training instances in memory
 - for each training instance $(x^{(i)}, y^{(i)})$
 - if other training instances provide correct classification for $(x^{(i)}, y^{(i)})$
 - delete it from the memory
- *incremental growth*
 - start with an empty memory
 - for each training instance $(x^{(i)}, y^{(i)})$
 - if other training instances in memory **don't** correctly classify $(x^{(i)}, y^{(i)})$
 - add it to the memory



k-d Tree: Data Structure for Finding Nearest Neighbors







k-d trees

a *k-d tree* is similar to a decision tree except that each internal node

- stores one instance
- splits on the median value of the feature having the highest variance





Finding nearest neighbors with a k-d tree

- use **branch-and-bound** search
- priority queue stores
 - nodes considered
 - lower bound on their distance to query instance
- lower bound given by distance using a **single** feature
- average case: $O(\log_2 m)$
- worst case: $O(m)$ where m is the size of the training-set


Finding nearest neighbors in a k-d tree



NearestNeighbor(instance $x^{(q)}$)


```
PQ = {}                                // minimizing priority queue
best_dist = ∞                            // smallest distance seen so far
PQ.push(root, 0)
while PQ is not empty
    (node, bound) = PQ.pop();
    if (bound ≥ best_dist)
        return best_node.instance          // nearest neighbor found
    dist = distance( $x^{(q)}$ , node. instance)
    if (dist < best_dist)
        best_dist = dist
        best_node = node
    if ( $q[\text{node.feature}] - \text{node.threshold} > 0$ )
        PQ.push(node.left,  $x^{(q)}[\text{node.feature}] - \text{node.threshold}$ )
        PQ.push(node.right, 0)
    else
        PQ.push(node.left, 0)
        PQ.push(node.right, node. threshold -  $x^{(q)}[\text{node.feature}]$ )
return best_node. instance
```

k-d tree example (Manhattan distance)




given query

$$\mathbf{x}^{(q)} = (2, 3)$$




k-d tree example (Manhattan distance)




given query

$$x^{(q)} = (2, 3)$$




distance	best distance	best node	priority queue
	∞		(f, 0)

k-d tree example (Manhattan distance)



given query


$$x^{(q)} = (2, 3)$$



pop f


distance	best distance	best node	priority queue
	∞		(f, 0)
4.0	4.0	f	

k-d tree example (Manhattan distance)



given query


$$\mathbf{x}^{(q)} = (2, 3)$$



pop f


distance	best distance	best node	priority queue
	∞		(f, 0)
4.0	4.0	f	(c, 0)

k-d tree example (Manhattan distance)



given query


$$x^{(q)} = (2, 3)$$



pop f


distance	best distance	best node	priority queue
	∞		(f, 0)
4.0	4.0	f	(c, 0) (h, 4)

k-d tree example (Manhattan distance)




given query

$$x^{(q)} = (2, 3)$$




	distance	best distance	best node	priority queue
pop f		∞		(f, 0)
pop c	4.0	4.0	f	(c, 0) (h, 4)
	10.0	4.0	f	

k-d tree example (Manhattan distance)




given query

$$x^{(q)} = (2, 3)$$




	distance	best distance	best node	priority queue
pop f		∞		(f, 0)
pop c	4.0	4.0	f	(c, 0) (h, 4)
	10.0	4.0	f	(e, 0) (h, 4) (b, 7)

k-d tree example (Manhattan distance)




given query

$$x^{(q)} = (2, 3)$$




	distance	best distance	best node	priority queue
pop f		∞		(f, 0)
pop c	4.0	4.0	f	(c, 0) (h, 4)
pop e	10.0	4.0	f	(e, 0) (h, 4) (b, 7)
	1.0	1.0	e	

k-d tree example (Manhattan distance)




given query

$$x^{(q)} = (2, 3)$$




	distance	best distance	best node	priority queue
		∞		(f, 0)
pop f	4.0	4.0	f	(c, 0) (h, 4)
pop c	10.0	4.0	f	(e, 0) (h, 4) (b, 7)
pop e	1.0	1.0	e	(d, 1) (h, 4) (b, 7)

k-d tree example (Manhattan distance)



given query

$$x^{(q)} = (2, 3)$$



	distance	best distance	best node	priority queue
pop f		∞		(f, 0)
pop c	4.0	4.0	f	(c, 0) (h, 4)
pop e	10.0	4.0	f	(e, 0) (h, 4) (b, 7)
pop d	1.0	1.0	e	(d, 1) (h, 4) (b, 7)
				return e

An aerial photograph of a city skyline at sunset. The city is built along a large body of water, with numerous buildings of various architectural styles and heights. A dense forest line runs along the water's edge. The sky is filled with warm, golden light from the setting sun. In the foreground, the dark blue water of the lake is dotted with many small sailboats and other watercraft.

Strength and Limitations





Strengths of instance-based learning

- simple to implement
- “training” is very efficient
- adapts well to on-line learning
- robust to noisy training data (when $k > 1$)
- often works well in practice



Limitations of instance-based learning

- sensitive to range of feature values
- sensitive to irrelevant and correlated features, although ...
 - there are variants (such as locally weighted regression) that learn weights for different features
 - later we'll talk about *feature selection* methods
- classification/prediction can be inefficient, although edited methods and k - d trees can help alleviate this weakness
- doesn't provide much insight into problem domain because there is no explicit model




Inductive bias

- *inductive bias* is the set of assumptions a learner uses to be able to predict y_i for a previously unseen instance x_i
- two components
 - *hypothesis space bias*: determines the models that can be represented
 - *preference bias*: specifies a preference ordering within the space of models
- in order to *generalize* (i.e. make predictions for previously unseen instances) a learning algorithm must have an inductive bias



Consider the inductive bias of DT and k -NN learners

learner	hypothesis space bias	preference bias
ID3 decision tree	trees with single-feature, axis-parallel splits	small trees identified by greedy search
k -NN	Voronoi decomposition determined by nearest neighbors	instances in neighborhood belong to same class



THANK YOU

Some of the slides in these lectures have been adapted/borrowed from materials developed by Yingyu Liang, Mark Craven, David Page, Jude Shavlik, Tom Mitchell, Nina Balcan, Elad Hazan, Tom Dietterich, and Pedro Domingos.

