

An aerial photograph of the University of Wisconsin-Madison campus during sunset. The sun is low on the horizon, casting a warm orange glow over the buildings and the surrounding greenery. The campus buildings are a mix of architectural styles, with many featuring red brick. A large, calm lake occupies the right side of the frame, dotted with numerous small sailboats and other watercraft. The overall atmosphere is peaceful and scenic.

Naïve Bayes

CS 760@UW-Madison





Goals for the lecture

- understand the concepts
 - generative/discriminative models
 - examples of the two approaches
 - MLE (Maximum Likelihood Estimation)
 - Naïve Bayes
 - Naïve Bayes assumption
 - model 1: Bernoulli Naïve Bayes
 - model 2: Multinomial Naïve Bayes
 - model 3: Gaussian Naïve Bayes
 - model 4: Multiclass Naïve Bayes



Review: supervised learning

problem setting

- set of possible instances: X
- unknown *target function* (concept): $f : X \rightarrow Y$
- set of *hypotheses* (hypothesis class): $H = \{h \mid h : X \rightarrow Y\}$

given

- *training set* of instances of unknown target function f

$$(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}) \dots (\mathbf{x}^{(m)}, y^{(m)})$$


output

- hypothesis $h \in H$ that best approximates target function



Parametric hypothesis class

- hypothesis $h \in H$ is indexed by (fixed dimensional) parameter $\theta \in \Theta$
- learning: find the θ such that $h_\theta \in H$ best approximate the target



- different from nonparametric approaches like decision trees and nearest neighbor
- advantages: various hypothesis class; easier to use math/optimization



Discriminative approaches

- hypothesis $h \in H$ directly predicts the label y given the features x

$$y = h(x) \text{ or more generally, } p(y | x) = h(x)$$

- then define a loss function $L(h)$ and find hypothesis with min. loss
 - A special case is a probabilistic model, finding MLE or MAP
- example: linear regression

$$h_{\theta}(x) = \langle x, \theta \rangle$$

$$L(h_{\theta}) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



Generative approaches

- hypothesis $h \in H$ specifies a **generative probabilistic story** for how the full data (x,y) was created

$$h(x, y) = p(x, y)$$

- then pick a hypothesis by maximum likelihood estimation (**MLE**) or Maximum A Posteriori (**MAP**)

- example: roll a weighted die
- weights for each side (θ) define how the data are generated
- use MLE on the training data to learn θ

Comments on discriminative/generative



- Orthogonal to the parametric / nonparametric divide
 - nonparametric Bayesian: a large subfield of ML
- when discriminative/generative is likely to be better? Discussed in later lecture
- typical discriminative: linear regression, logistic regression, SVM, many neural networks (not all!), ...
- typical generative: Naïve Bayes, Bayesian Networks, ...

An aerial photograph of a city skyline at sunset, likely Madison, Wisconsin. The city is built along a riverbank, with numerous buildings of varying architectural styles and a dense forest area behind them. The river is filled with many small sailboats and other watercraft. The sky is a warm, golden color from the setting sun.

MLE and MAP





MLE vs. MAP

Suppose we have data $\mathcal{D} = \{x^{(i)}\}_{i=1}^N$

$$\theta^{\text{MLE}} = \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^N p(\mathbf{x}^{(i)} | \boldsymbol{\theta})$$

Maximum Likelihood
Estimate (MLE)



Background: MLE

Example: MLE of Exponential Distribution

- pdf of $\text{Exponential}(\lambda)$: $f(x) = \lambda e^{-\lambda x}$
- Suppose $X_i \sim \text{Exponential}(\lambda)$ for $1 \leq i \leq N$.
- Find MLE for data $\mathcal{D} = \{x^{(i)}\}_{i=1}^N$
- First write down log-likelihood of sample.
- Compute first derivative, set to zero, solve for λ .
- Compute second derivative and check that it is concave down at λ^{MLE} .



Background: MLE

Example: MLE of Exponential Distribution

- First write down log-likelihood of sample.

$$\ell(\lambda) = \sum_{i=1}^N \log f(x^{(i)}) \quad (1)$$

$$= \sum_{i=1}^N \log(\lambda \exp(-\lambda x^{(i)})) \quad (2)$$

$$= \sum_{i=1}^N \log(\lambda) + -\lambda x^{(i)} \quad (3)$$

$$= N \log(\lambda) - \lambda \sum_{i=1}^N x^{(i)} \quad (4)$$



Background: MLE

Example: MLE of Exponential Distribution

- Compute first derivative, set to zero, solve for λ .

$$\frac{d\ell(\lambda)}{d\lambda} = \frac{d}{d\lambda} N \log(\lambda) - \lambda \sum_{i=1}^N x^{(i)} \quad (1)$$

$$= \frac{N}{\lambda} - \sum_{i=1}^N x^{(i)} = 0 \quad (2)$$

$$\Rightarrow \lambda^{\text{MLE}} = \frac{N}{\sum_{i=1}^N x^{(i)}} \quad (3)$$



MLE vs. MAP

Suppose we have data $\mathcal{D} = \{x^{(i)}\}_{i=1}^N$

$$\theta^{\text{MLE}} = \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^N p(\mathbf{x}^{(i)} | \boldsymbol{\theta})$$

Maximum Likelihood Estimate (MLE)

$$\theta^{\text{MAP}} = \underset{\theta}{\operatorname{argmax}} \prod_{i=1}^N p(\mathbf{x}^{(i)} | \boldsymbol{\theta}) p(\boldsymbol{\theta})$$

Maximum *a posteriori* (MAP) estimate



Prior

An aerial photograph of a city skyline at sunset, likely Madison, Wisconsin. The city is built along a large body of water, with numerous sailboats and small boats scattered across the surface. The sky is filled with warm, golden light from the setting sun. The city buildings, including several prominent red brick structures, are visible along the shoreline.

Naïve Bayes



Model 0: Not-so-naïve Model?



Generative Story:

1. Flip a weighted coin (Y)
2. If heads, roll the **yellow** many sided die to sample a document vector (X) from the Spam distribution
3. If tails, roll the **blue** many sided die to sample a document vector (X) from the Not-Spam distribution

$$P(X_1, \dots, X_K, Y) = P(X_1, \dots, X_K | Y)P(Y)$$

This model is
computationally naïve!





Model 0: Not-so-naïve Model?

Generative Story:

1. Flip a weighted coin (Y)
2. If heads, sample a document ID (X) from the Spam distribution
3. If tails, sample a document ID (X) from the Not-Spam distribution

$$P(X, Y) = P(X|Y)P(Y)$$


This model is
computationally naïve!






Model 0: Not-so-naïve Model?

Flip weighted coin




If HEADS, roll
yellow die



Each side of the die
is labeled with a
document vector
(e.g. [1,0,1,...,1])

If TAILS, roll
blue die



y	x_1	x_2	x_3	...	x_K
0	1	0	1	...	1
1	0	1	0	...	1
1	1	1	1	...	1
0	0	0	1	...	1
0	1	0	1	...	0
1	1	0	1	...	0

Naïve Bayes Assumption



Conditional independence of features:

$$\begin{aligned} P(X_1, \dots, X_K, Y) &= P(X_1, \dots, X_K | Y)P(Y) \\ &= \left(\prod_{k=1}^K P(X_k | Y) \right) P(Y) \end{aligned}$$



C	P(C)
0	0.33
1	0.67

Estimating a joint from conditional probabilities

$$P(A, B | C) = P(A | C) * P(B | C)$$

A	C	P(A C)
0	0	0.2
0	1	0.5
1	0	0.8
1	1	0.5

A	B	C	P(A,B,C)
0	0	0	...
0	0	1	...
0	1	0	...
0	1	1	
1	0	0	
1	0	1	
1	1	0	
1	1	1	

B	C	P(B C)
0	0	0.1
0	1	0.9
1	0	0.9
1	1	0.1



C	P(C)
0	0.33
1	0.67

A	C	P(A C)		
0	0	0.2		
0	1	0.5		
1	0	0.8		
1	B	C	P(B C)	
0	0	0.1		
0	1	0.9		
1	0	0.9		
1	1	0.1		

Estimating a joint from conditional probabilities

A	B	D	C	P(A,B,D,C)
0	0	0	0	...
0	0	1	0	...
0	1	0	0	...
0	1	1	0	
1	0	0	0	
1	0	1	0	
1	1	0	0	
1	1	1	0	
...

D	C	P(D C)
0	0	0.1
0	1	0.1
1	0	0.9
1	1	0.1



Assuming conditional independence, the conditional probabilities encode the **same information** as the joint table.

They are very convenient for estimating

$$P(X_1, \dots, X_n | Y) = P(X_1 | Y) * \dots * P(X_n | Y)$$

They are almost as good for computing

$$P(Y | X_1, \dots, X_n) = \frac{P(X_1, \dots, X_n | Y)P(Y)}{P(X_1, \dots, X_n)}$$

$$\forall \mathbf{x}, y : P(Y = y | X_1, \dots, X_n = \mathbf{x}) = \frac{P(X_1, \dots, X_n = \mathbf{x} | Y)P(Y = y)}{P(X_1, \dots, X_n = \mathbf{x})}$$

Generic Naïve Bayes Model



Support: Depends on the choice of **event model**, $P(X_k|Y)$

Model: Product of **prior** and the event model

$$P(\mathbf{X}, Y) = P(Y) \prod_{k=1}^K P(X_k|Y)$$

Training: Find the **class-conditional** MLE parameters

For $P(Y)$, we find the MLE using all the data. For each $P(X_k|Y)$ we condition on the data with the corresponding

Classification: Find the class that maximizes the posterior

$$\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x})$$

Generic Naïve Bayes Model



Classification:

$$\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x}) \quad (\text{posterior})$$

$$= \operatorname{argmax}_y \frac{p(\mathbf{x}|y)p(y)}{p(x)} \quad (\text{by Bayes' rule})$$

$$= \operatorname{argmax}_y p(\mathbf{x}|y)p(y)$$



Various Naïve Bayes Models





Model 1: Bernoulli Naïve Bayes

Support: Binary vectors of length K

$$\mathbf{x} \in \{0, 1\}^K$$

Generative Story:

$$Y \sim \text{Bernoulli}(\phi)$$

$$X_k \sim \text{Bernoulli}(\theta_{k,Y}) \quad \forall k \in \{1, \dots, K\}$$

Model: $p_{\phi, \boldsymbol{\theta}}(\mathbf{x}, y) = p_{\phi, \boldsymbol{\theta}}(x_1, \dots, x_K, y)$


$$= p_\phi(y) \prod_{k=1}^K p_{\boldsymbol{\theta}_k}(x_k | y)$$

$$= (\phi)^y (1 - \phi)^{(1-y)} \prod_{k=1}^K (\theta_{k,y})^{x_k} (1 - \theta_{k,y})^{(1-x_k)}$$



Model 1: Bernoulli Naïve Bayes

Flip weighted coin



If HEADS, flip each yellow coin



If TAILS, flip each blue coin



y	x_1	x_2	x_3	\dots	x_K
0	1	0	1	\dots	1
1	0	1	0	\dots	1
1	1	1	1	\dots	1
0	0	0	1	\dots	1
0	1	0	1	\dots	0
1	1	0	1	\dots	0

Each red coin corresponds to an x_k

We can **generate** data in this fashion. Though in practice we never would since our data is **given**.

Instead, this provides an explanation of **how** the data was generated (albeit a terrible one).



Model 1: Bernoulli Naïve Bayes

Support: Binary vectors of length K

$$\mathbf{x} \in \{0, 1\}^K$$

Generative Story:

$$Y \sim \text{Bernoulli}(\phi)$$

$$X_k \sim \text{Bernoulli}(\theta_{k,Y}) \quad \forall k \in \{1, \dots, K\}$$

Model: $p_{\phi, \theta}(x, y) = (\phi)^y (1 - \phi)^{(1-y)} \prod_{k=1}^K \theta_{k,y}^{x_k} (1 - \theta_{k,y})^{1-x_k}$

Same as Generic
Naïve Bayes

Classification: Find the class that maximizes the posterior

$$\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x})$$

Generic Naïve Bayes Model

Recall... 

Classification:

$$\hat{y} = \operatorname{argmax}_y p(y|\mathbf{x}) \quad (\text{posterior})$$

$$= \operatorname{argmax}_y \frac{p(\mathbf{x}|y)p(y)}{p(x)} \quad (\text{by Bayes' rule})$$

$$= \operatorname{argmax}_y p(\mathbf{x}|y)p(y)$$



Model 1: Bernoulli Naïve Bayes

Training: Find the **class-conditional** MLE parameters

For $P(Y)$, we find the MLE using all the data. For each $P(X_k|Y)$ we condition on the data with the corresponding class.

$$\phi = \frac{\sum_{i=1}^N \mathbb{I}(y^{(i)} = 1)}{N}$$

$$\theta_{k,0} = \frac{\sum_{i=1}^N \mathbb{I}(y^{(i)} = 0 \wedge x_k^{(i)} = 1)}{\sum_{i=1}^N \mathbb{I}(y^{(i)} = 0)}$$

$$\theta_{k,1} = \frac{\sum_{i=1}^N \mathbb{I}(y^{(i)} = 1 \wedge x_k^{(i)} = 1)}{\sum_{i=1}^N \mathbb{I}(y^{(i)} = 1)}$$

$$\forall k \in \{1, \dots, K\}$$



Model 2: Multinomial Naïve Bayes

Support:

Integer vector (word IDs)

$\mathbf{x} = [x_1, x_2, \dots, x_M]$ where $x_m \in \{1, \dots, K\}$ a word id.

Generative Story:

for $i \in \{1, \dots, N\}$:

$y^{(i)} \sim \text{Bernoulli}(\phi)$

for $j \in \{1, \dots, M_i\}$: (Assume $M_i = M$ for all i)

$x_j^{(i)} \sim \text{Multinomial}(\boldsymbol{\theta}_{y^{(i)}}, 1)$

Model:

$$\begin{aligned} p_{\phi, \boldsymbol{\theta}}(\mathbf{x}, y) &= p_{\phi}(y) \prod_{k=1}^K p_{\boldsymbol{\theta}_k}(x_k | y) \\ &= (\phi)^y (1 - \phi)^{(1-y)} \prod_{j=1}^{M_i} \theta_{y, x_j} \end{aligned}$$



Model 3: Gaussian Naïve Bayes

Support:

$$\mathbf{x} \in \mathbb{R}^K$$

Model: Product of **prior** and the event model

$$p(\mathbf{x}, y) = p(x_1, \dots, x_K, y)$$

$$= p(y) \prod_{k=1}^K p(x_k|y)$$

Gaussian Naive Bayes assumes that $p(x_k|y)$ is given by a Normal distribution.



Model 4: Multiclass Naïve Bayes


Model:

The only change is that we permit y to range over C classes.

$$p(\mathbf{x}, y) = p(x_1, \dots, x_K, y)$$

$$= p(y) \prod_{k=1}^K p(x_k|y)$$

Now, $y \sim \text{Multinomial}(\phi, 1)$ and we have a separate conditional distribution $p(x_k|y)$ for each of the C classes.



THANK YOU

Some of the slides in these lectures have been adapted/borrowed from materials developed by Yingyu Liang, Mark Craven, David Page, Jude Shavlik, Tom Mitchell, Nina Balcan, Matt Gormley, Elad Hazan, Tom Dietterich, and Pedro Domingos.

